

Ripple Protocol Consensus Algorithm Review

Peter Todd

May 11th 2015

1 Activities Performed

Four major activities were performed during the review:

1. Reviewed white papers and development documentation at <https://ripple.com>.
2. Reviewed existing third-party criticism.
3. Setup and run a full node.
4. Reviewed C++ implementation, specifically the 0.27.4 tag of `rippled`.¹

2 Overall Architecture

Ripple has diverged significantly from the original concept[3] of a decentralized network recording the transferral of money explicitly as debt relationships between parties. (Figure 1, showing entities A, B, and C participating in a transaction) While the concept of a debt relationship still exists in the form of trustlines between participants, the bulk of the codebase and developer documentation now focuses² on the use and maintenance of a global ledger of transactions and account balances. Additionally a native currency, XRP, has been added, which is used for anti-spam transaction fees³ and to serve as a universal currency. (Figure 2)

3 Global Ledger

Figure 3 shows the basic cryptographic structure of the Ripple ledger. As with Bitcoin, it consists of a set of chain of blocks, each consisting of a block header

¹`git commit 92812fe7239ffa3ba91649b2ece1e892b866ec2a` from <https://github.com/ripple/rippled>

²For instance the “Tutorials” section of the Ripple website only explains how to create transactions that modify the ledger; there is almost no information available on how to actually use the trustlines feature.

³A small amount of XRP is irrevocably destroyed for every transaction.




Figure 1: Original Ripple Architecture




Figure 2: Ripple Labs Architecture

that commits to a previous block and a set of new transactions for that block. (via a merkle tree) Additionally each block header also commits to a merkelized trie⁴ of the state of all account balances. The block chain doesn't itself deal directly with the consensus protocol. Rather validators sign messages as part of the consensus protocol, and knowledge of those messages lets Ripple nodes determine the correct chain.

Detailed documentation about exactly how the ledger is structured is spotty; the descriptions here are mostly based on direct analysis of the `rippled` source code.

⁴Oddly the trie is not binary, trading off performance for significantly larger proofs.




Figure 3: Ripple Ledger Structure

3.1 Transactions

Unlike Bitcoin, transactions increment and decrement account balances; they do not directly consume the output of other transactions. Transactions are also used to setup and modify trustlines. To prevent replay attacks, accounts and transactions have sequence numbers; a transaction is only valid if the Sequence number is exactly one greater than the last-validated transaction from the same account. Additionally an optional `AccountTxnID` field is available for use within transactions; the transaction is only valid if the sending account’s previously-sent transaction matches the provided hash.

Authentication is performed with a simple signature scheme; a scripting system is not available, nor is multi-signature support. To implement the latter a highly complex single-purpose scheme[7] has been proposed.

3.2 Low-Level Hashing and Serialization

Though various parts of the Ripple user API’s and networking use industry standard serialization formats like JSON and Google Protobuf, for consensus-critical functionality Ripple uses a custom tag-length-value serialization and hashing scheme. This is not unexpected as industry standard serialization schemes rarely, if ever, account for the need to create hash digests from represented objects.

$$\text{Serialize}(\text{obj}) = t_0 + n_0 + d_0 + \dots + t_n + n_n + d_n \quad (1)$$

Hashing of objects generally uses the following scheme, implemented by `STObject::getHash()`, resulting in a 256bit digest:

$$H(\text{obj}) = \text{SHA512}(p + \text{Serialize}(\text{obj}))[0 : 256\text{bits}] \quad (2)$$

The prefix p is per-object-type, guaranteeing that objects of different types will always have a different hash.⁵ For objects containing signatures, such as transactions, there is a similar but separate *signing hash* implemented by the function `STObject::getSigningHash()`. Unlike the standard hash, the signing hash does not serialize signature fields.

3.3 Open Questions

- How exactly does the block header skiplist work? While the source code refers to a skiplist, the actual block header format appears to only allow block headers to refer to a single previous block header in a linear fashion.
- The `rippled` codebase suggests that in some circumstances account balances can be negative. What are those circumstances exactly?
- Do opportunities exist to “shard” the Ripple blockchain to improve scaling by distributing load across parallel servers?

⁵This is commonly known as *tagged hashing* in the literature.

- What plans exist (if any) to support SPV-style “lite clients” in the Ripple protocol?

4 Consensus Algorithm

Coming to consensus about the contents of the Ripple ledger is the role of the Ripple Protocol Consensus Algorithm[5]. In situations where the participants can be determined in advance, consensus algorithms are a well studied with a large variety of choices available. With that in mind while this review did some basic “sanity checking” of the specific algorithm, for the most part we assume it works correctly as described. Instead we’ll focus on the assumptions surrounding the consensus algorithm, and how it can and will be used in practice.

4.1 Unique Node List

Like the vast majority of consensus algorithms, the Ripple consensus algorithm starts with a known set of nodes known to be participating in the consensus. This Unique Node List, (UNL) is a list of public keys meant to be associated with active (validating) nodes the node operator believes are “unique”. [8] Ripple Labs suggests that UNL’s “should have 100+ nodes on them.” [8] and provides a “starter” UNL at <https://ripple.com/ripple.txt> with 5 to 8 nodes.⁶

Through the consensus algorithm nodes on the UNL vote to determine the contents of the Ripple ledger. While the actual protocol contains a number of rounds of proposals and voting the end result can be described as basically a supermajority vote: a transaction is only approved if 80% of the UNL of a server agrees with it.[5, 3.2] Put another way, if 20% of the UNL choose to reject a transaction, it will not be included in the ledger.

Different nodes may chose to use different UNLs; if the UNLs do not sufficiently overlap global consensus is not guaranteed because different UNL “cliques” can come to consensus independently of each other. For instance Figure 4 shows a simple example with two UNL cliques, red and blue, that have forked because of insufficient connectivity, resulting in the forked blockchain shown in Figure 5. More complex failure modes have also been identified, such as the “less than majority evil” failure mode identified by Gregory Maxwell,[10] and “separate supermajorities” mode identified by Andrew Miller.[11]

Increasing the connectivity past the 20% threshold, as shown in Figure 6, brings both cliques into agreement. However what happens to the blockchain? The Ripple Protocol Consensus Algorithm paper is unclear on this point. For instance, the loss and restoration of connectivity may have been associated with network latency. The paper does state that “nodes whose latency grows larger than a preset bound b are removed from all UNLs”[5, 3.4.1] but does not say

⁶Exactly how this UNL is generated is unknown; the author downloaded it on multiple occasions getting between 5 to 8 nodes each time, mostly the same.




Figure 4: Disjoint UNL




Figure 5: Forked blockchain

if this removal is meant to be permanent⁷ nor if by “all” they refer to removal from UNLs on all nodes in the network. Additionally while the author did not have sufficient time to investigate this point fully, the source code doesn’t appear to have any provision for reorganizations of the ledger when a longer fork is detected.

Regardless, it is certain that following a fork caused by disjoint UNL cliques the losing side simply has to discard some or all history (Figure 7) possibly resulting in double-spends and associated financial losses.

⁷In the existing implementation restarting the node does reload the UNL set from the config file.




Figure 6: UNL cliques in consensus




Figure 7: Post-fork reorganization

		Majority		
		100%	50%	10%
Local	100%	Consensus	DoS	DoS
	50%	DoS	DoS	Double-spend
	10%	Double-spend	Double-spend	Double-spend

Table 1: Worst-case outcomes of default UNL % choices

4.2 Choosing the Unique Node List

How should Ripple node administrators pick their Unique Node List? Ripple Labs provides little, if any, concrete advice beyond the “starter UNL” they provide. Let’s look at this problem from another angle: Why shouldn’t a node operator just use the default UNL?

Let’s assume the node operator does not know what UNL other nodes in the economic majority have chosen to use. They do however know the default UNL provided by Ripple Labs, and they have a reasonable belief that other node operators also know that UNL and may also be using it. The node operator’s goal is to limit losses; the best case is their node stays in consensus with other nodes. Second best is they suffer a denial of service attack causing their node to stop processing some or all transactions. Worst case is their node accepts transactions that other nodes do not, making possible double-spend attacks. Table 1 shows the outcomes of various possible decisions by them and other node operators, ranging from 100% default UNL, 50% default UNL, and 10% default UNL.

In every circumstance node operators can reduce the risk of losses due to consensus failures by removing non-default UNL entries, with the least risky option being to simply stick with the default UNL of 100% Ripple Labs controlled nodes. This is particularly acute given the low profit margins of many financial transactions: while a denial-of-service halts incoming revenue, a double-spend attack can quickly drain working capital, quite possibly without any hope of ever recovering the stolen funds.

A closely related issue is what incentive does a node operator have to publicly perform validation services? Ripple currently has no compensation mechanism for public validators, yet validation at minimum raises potential legal issues such as lawsuits for negligence or aiding financial crimes. Again, the least-risk option is to not publicly validate the Ripple ledger.

5 Attack Scenarios

To evaluate potential attacks on the Ripple network we look at three factors:

- **Type:** What is the effect of the attack? Denial of Service and/or theft? Or is the attack mostly useful in preparation for another type of attack?
- **Cost:** What is a probable minimum cost to carry out the attack, for the people able to do the attack? For instance, a technical attack that can be launched by “bored teenagers” will have an essentially zero cost while a long-term rewrite of a proof-of-work blockchain has a reasonably well-defined cost.
- **Scope:** Who is affected by the attack? Is this a targeted attack, mainly affecting a one or two targets? An attack with broad impact? Or an attack with global impact on all Ripple users?
- **Duration:** How long will it take before the attack is neutralized? This may be a matter of hours to write and distribute a simple software patch⁸, to indefinitely long for some attacks where recovery is impossible.
- **Probability:** How likely is it for the Ripple network to be attacked this way?

We do *not* try to determine the monetary impact of the attack, because we are unable to predict how the Ripple network will be used. For instance, while a double-spend attack may have the same basic impact on a \$100/day small business as it would on a \$1,000,000,000/day multinational, the latter could lose orders of magnitude more money for what it at a technical level basically the same attack. A closely issue with attempting to predict monetary losses is that at some point losses can be sufficiently high that they create situations where social consensus can decide that the losses simply must be reversed. For instance, Vericoin community chose to do a hard-fork of their alt-currency to undo the impact of a major theft.[6]

5.1 Consensus Split

Type	Cost	Scope	Duration	Probability
DoS	\$0	Global	Hours	Very High
Theft	\$1k	Targeted	Hours	Medium

The attacker exploits a difference in behavior between different implementations/versions of the Ripple protocol. The result of the attack can be a simple denial of service due to the Ripple network being unable to process transactions, or with more sophistication, the attacker can fork the Ripple ledger/blockchain and use the fork to double-spend. The attack is stopped by identifying it, and

⁸The author personally stopped the CVE-2013-4627 attack on the Bitcoin network in a few hours of rushed work in the middle of the night.

then distributing software patches. Provided the Ripple network is well monitored consensus split attacks can be stopped fairly quickly; the March 2013 Bitcoin chain fork[1] was resolved in a matter of hours by contacting mining pools. A similar response could happen on Ripple by contacting validating nodes commonly used in UNLs.⁹

Note that a consensus split may also happen by accident; prior consensus splits in cryptocurrency systems have almost always been accidental. Equally, that consensus splits can happen by accident clearly shows that the minimum cost to perform this attack is zero. However, due to the limited number of consensus bugs available to exploit in a given (set of) versions of the Ripple protocol in use and the fast response time possible to consensus splits the potential of this class of attack for long-term attacks on the network is limited.

Risk factors:

- The main implementation of the Ripple protocol, `rippled`, does not cleanly separate the consensus-critical part of the codebase from the non-consensus-critical part.¹⁰
- The Ripple protocol itself is extremely complex, with many different types of transactions, and all functionality being implemented directly in the protocol itself.
- Improvements to Ripple frequently require changing the consensus-critical codebase rather than user software. For instance while on Bitcoin the implementation of multisig was possible without modification to the protocol[4] in Ripple the lack of extension capabilities such as scripting require a consensus-critical change.[7]¹¹

5.2 Transaction Flood

Type	Cost	Scope	Duration	Probability
DoS	> \$100k	Broad	Varies	Low

The attacker creates large numbers of transactions - either covertly or overtly - until the Ripple network fails due to overload, or more likely, Ripple users are priced out of their ability to use the Ripple network. Transactions pay a minimum fee per transaction by destroying XRP; the fee per transaction is set by a vote between all validators on the UNL list. In the short to medium term such an attack would quickly drive up that fee. A successful attack requires deep pockets as the attacker must out-spend other Ripple users.

Risk factors:

- Ripple Labs plans[14] to require verified user identification to use the Ripple network in the near future. When implemented this will make

⁹ Another reason to use the UNL provided Ripple Labs.

¹⁰ The consensus-critical portions of Bitcoin Core codebase are currently being separated into a stand-alone `libconsensus` library.

¹¹ P2SH was *not* required for multisig to be implemented.

it easy to determine who is doing a covert transaction flood attack and restrict their access to the network.

- The global consensus ledger has inherently poor $O(n^2)$ scaling; even in the absence of a malicious attack the poor scalability of the ledger is a threat to the viability of Ripple.
- A clever attacker wishing to overtly attack the Ripple network can masquerade their traffic as legit economic activity. From the point of view of Ripple Labs - a major XRP owner - such an attack may not even be considered an attack!¹²
- Attacking the poor scalability of Ripple is highly synergistic with attacks depending on the high centralization of the Ripple network as even “un-successful” scalability attacks drive further centralization. For instance, if the response to a scalability attack is for Ripple Labs to beef up the validation nodes they control, and users stop running full nodes themselves, the network is more vulnerable to any attacks on those centrally managed validation nodes.

5.3 Coercion of Validators

Type	Cost	Scope	Duration	Probability
DoS	> \$100k	Targeted	Indefinite	High
DoS	> \$1M	Global	Indefinite	Medium
Theft	> \$100k	Targeted	Indefinite	Medium
Theft	> \$1M	Global	Indefinite	Low

Validators are coerced through political/legal/criminal means into attacking some or all Ripple users by preventing valid transactions from entering the ledger (DoS) and/or by allowing invalid transactions to enter the ledger. (theft) The cost estimates above are very rough, and assuming the attack is via legal means; the cost may potentially be as low as zero in some scenarios like blackmail. In preparation for this attack the attackers may also attempt to reduce the number of validators by with, for instance, a transaction flood attack.

Risk factors:

- The default UNL list and incentives against operating or using non-official validators discussed in Section 4.2 significantly increase the effectiveness of a coercion attack as nearly all Ripple users will be affected. Equally, in the event of a coercion attack there are still significant downsides to adopting a different UNL.
- Similarly this attack is synergistic with transaction flooding and other attacks on the scalability of the network, reducing the number of validators that need to be coerced.

¹²Similar to the debates in the Bitcoin ecosystem about whether or not fee-paying uses of the Bitcoin blockchain for non-Bitcoin-denominated transactions constitute “attacks”.

- The **rippled** node implementation does not check full history. This means if validators sign an invalid ledger, the transactions that make it invalid can be hidden from Ripple users. Equally, even with detection, the fact that Ripple users can be asked to simply restart/reinstall their nodes to get past the invalid history may significantly reduce the social barriers to creating an invalid ledger.
- The Ripple protocol is constructed such that it is possible to create many, although not all, types of fraud proofs[9, 15] required to prove an invalid ledger. While the software does not currently implement fraud proofs, this is functionality that could be added in the future without a change to the consensus. If widely used fraud proofs could prevent some types of coercion attacks that force validators to create invalid ledgers by making it likely for the result of the attack to be DoS rather than profitable theft. In short, the attack would simply shut the Ripple network down for all users rather than result in a successful theft.
- Some people would argue that the recently announced plans[14] to require verified user identification to use the Ripple network in the near future are an example of a successful coercion attack. It can be argued that Ripple Labs is increasing the risk of coercion attacks by simultaneously adopting a “regulator friendly” position and promoting the use of Ripple across many different jurisdictions.

5.4 Software Backdoor

Type	Cost	Scope	Duration	Probability
Varies	\$0	Broad	Weeks	High

The attacker inserts a backdoor into a widely used Ripple protocol implementation. (e.g. the **rippled** node software) The attacker may be an internal Ripple Labs employee, an external member of the public submitting a pull-request, or they may compromise the software actually downloaded by end-users by compromising binary or source code hosting. (e.g. GitHub)

Risk factors:

- All relevant software is open source, allowing for third-party review. Ripple Labs also has a bug bounty program, which can incentivise that review.
- Ripple Labs appears to use GitHub - a third party - as their master source code repository, with development happening through pull-requests on GitHub. While this does place trust in a third-party, PGP signing source code during development can mitigate this risk. However there appears to be no PGP signatures on any git commits, giving Ripple Labs little ability to audit who wrote what code.
- **Currently Ripple Labs does not provide a secure way to download any of their software.** While binaries are available, the mechanism

to download them is a Ubuntu package repository over insecure HTTP; while source code is available through git, neither commits nor release tags are signed in anyway. **This is a serious omission that has lead to significant monetary losses in the past.** Ripple Labs should be following industry best-practice by signing git commits and tags[16] as well as PGP signing their Ubuntu packages.

5.5 Theft of Validator Secret Keys

Type	Cost	Scope	Duration	Probability
Precursor	\$0	N/A	N/A	Unknown

The attacker steals illegitimate control of the secret keys used by validators, particularly the validators in the default UNL. The theft may happen via an exploit, backdoor, compromised hosting company, etc. While not useful for the attacker in of itself, the stolen keys can be used for further attacks such as DoS or a simulated ledger attack.

A closely related attack is to compromise the Ripple website, and trick people into downloading a default UNL list consisting of only attacker controlled nodes.

Risk factors:

- Due to the high centralization of UNL validators, a relatively small number of validators need to be compromised for a successful attack.
- Most validators are running nearly the same, if not the exact same, software, making it likely for backdoor/exploit attack to successfully exfiltrate the secret keys. Due to the high consensus-critical complexity of the Ripple protocol and lack of a consensus library alternate validator implementations are currently infeasible, making it difficult to change this situation.
- The `rippled` implementation doesn't yet appear to support separating the validation/consensus and signing functionality into different machines, or even different user accounts. (possibly Ripple Labs has internal-use-only improvements)
- Details about the operational security practices of the validators on the default UNL is unavailable, making it difficult to assess risk.
- The Ripple P2P protocol does make it possible to operate validators without the nodes themselves accepting public traffic; the signatures produced by the validating nodes are distributed via an anonymizing flood-fill network. This setup is recommended by Ripple Labs documentation and makes remote attacks significantly more difficult.

5.6 Simulated Ledger

Type	Cost	Scope	Duration	Probability
Theft	> \$1k	Targeted	Days	N/A

An attacker with the ability to create signatures from nodes on the UNL can create an entirely simulated ledger in parallel to the “official” Ripple ledger, containing any transaction the attacker wishes. In conjunction with sybil attacking the target the attacker can then perform double-spend attacks. Sybil attacks are relatively cheap, requiring just the acquisition of a large number of IP addresses, e.g. via cloud hosting and/or botnet rental. In some situations sybil attacks can even have zero cost, such as the internal IT staff within the target with access to network infrastructure, the target’s ISP, or government actors with MITM capabilities such as the NSA.[13]

Risk factors:

- Because the Ripple consensus is determined by signatures creating a simulated ledger is a zero-cost attack, the “costless simulation”[12, 4.2] problem. Proof-of-work ledgers in comparison have very well-defined costs that discourage attacks and incentivise detection.¹³
- There are no direct incentives built into the Ripple protocol to discourage validators from allowing their private keys to sign simulated ledgers. By comparison the Slasher proof-of-stake algorithm attempts to disincentivise simulated history by destroying valuable bonds owned up by participants in the consensus process.[2]
- The lack of incentive may invite coercion attacks, as the validators being coerced into signing a simulated ledger have no direct incentive not too. For instance a court could order a validator to assist in a simulation attack on a target as a way of recovering funds held by the target.
- Simulation attacks are made significantly easier by the fact that the Ripple consensus protocol will remove validators from the UNL if they appear to have failed.[5, 3.4.2] With the target sybil attacked the attacker can prevent the target from learning that the other validators have *not* failed. When those validators are removed from the UNL due to the simulated failure the attacker does not need to compromise/coerce as many validators as before. Equally, if the target is using network split detection, the attacker can make the non-compromised validators *actually* fail as preparation for the attack. Again, the lack of incentive to run a validator is problematic.

¹³On the Bitcoin network even if the attacker can steal the hashing power needed to simulate history for free the miners involved lose money from the attack and have strong monetary incentives to detect and stop it.

6 Attack Flowchart

Figure 8 shows some of the stronger synergies between different attacks.

Let's go through a full scenario involving multiple attacks at once. Ripple Labs has diversified their default UNL such that the 80% of the validators are run by publicly known institutions evenly spread across the following five countries, United States, France, Switzerland, Russia, and China, with the remaining fifth being controlled by pseudoanonymous "DarkWeb" entities like The Silk Road Five, and Agora Market. As any one entity only comprises at most 1/6th of the UNL no single actor can block transactions.

The Russian government has issued a court order seizing assets of Royal Dutch Shell, who has recently received funds from Naftogaz Ukraine to develop a new fracking project. Naftogaz in turn owes Russia for unpaid natural gas imports - quite a lot more than the funds Shell has, but simply punishing them for working with Naftogaz is a good start.

By itself Russia is unable to freeze Shell's assets on the Ripple ledger through coercion of Russian-jurisdiction validators. So Russia begins with a transaction flood attack, slowly increasing load on the Ripple network to 10MB/sec of transactions. As a cover story a new gambling-oriented MMORPG is created that uses the Ripple ledger to record user accounts; Ripple Labs is happy to sell them the XRP to cover transaction fees.

The transaction flood results in the pseudoanonymous "DarkWeb" validators quitting - the Tor network is too slow to handle the load, and they worry they're real identities will be discovered if they move their validator nodes to clearnet hosting.

At this point Russia's 20% UNL control could simply vote to freeze Shell's assets, achieving a denial-of-service. However, it's decided it would be better if the funds were actually recovered instead.

Russia sets up a fake Chinese oil importer, hiring another Chinese company to do the actual transportation, and has them contact Shell asking to purchase some oil. They also find a wealthy Chinese investor with \$100 million USD worth of yuan in an account with a major Chinese bank, and has that investor "invest" in the new importer. Meanwhile Russia's cyberwarfare team gets to work, successfully compromising the Swiss and French validation nodes, and using the fact that they have agents in major Ukrainian ISPs to sybil attack the Ukrainian-based Shell office actually handling the sale. They can't compromise the US nodes - they're on a newer version - but they do happen to find a new zero-day consensus bug in the newer version.

On the day of the sale unknown to Shell their Ukrainian office's internet connection is being filtered. The IT staff at their accounting office voices concern when their Ripple node monitoring tools say that the US validators just went off-line, although they're somewhat reassured when they hear news that the US nodes just hit a consensus issue. The captain says there's bad weather coming in, so they need to leave the port ASAP - with France and Switzerland approving the transaction they have confidence that the Free World has approved the money transfer.

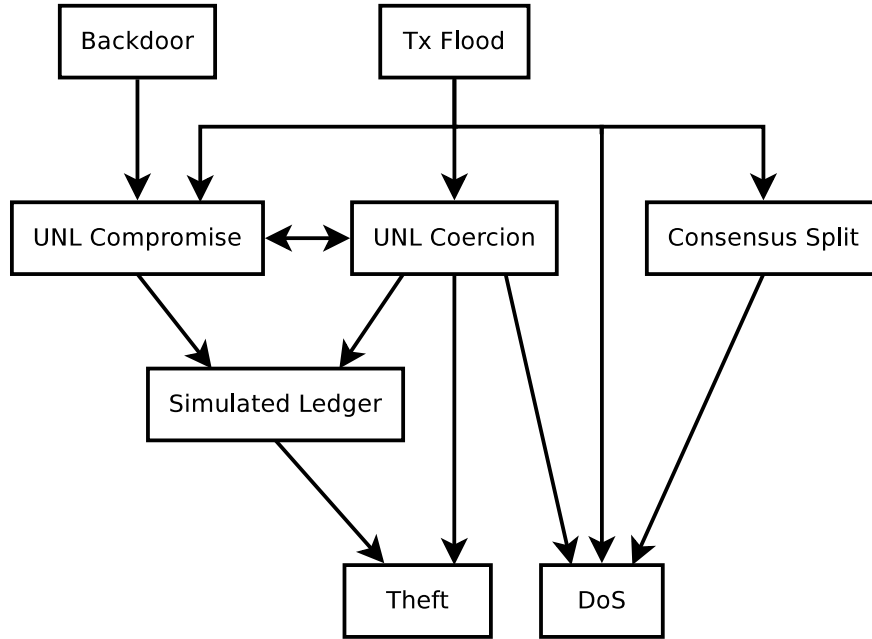


Figure 8: Attack flowchart

What has actually happened is Russian cyberwarfare experts created a fake consensus using stolen Swiss and French private keys. Russia of course officially signed this fake consensus, as well as the real consensus, taking the position this was an above-the-board court ordered operation; China also signed, but doesn't officially say why. The wealthy Chinese investor conveniently claims to have suddenly found out about the “fake” oil importer and on the real Ripple ledger has the funds frozen by the Chinese bank backing them - not a big deal with the money still sitting in the original account he deposited it too, still under Chinese jurisdiction.

Shell's Ukrainian office finally realises something is wrong after two hours of debugging, trying to figure out why an exchange wasn't seeing their transaction attempting to sell that yuan for USD. By that time the ship is in international waters, and is happy to keep steaming to China, where upon the oil is kept for the next ten years while the Chinese courts argue what to do with it.

7 Conclusions and Future Work

Nearly all the attack scenarios analysed above are direct outcomes of the need for consensus; the actual blockchain technology that records the ledger is relatively uninteresting. That the consensus requirement is for global consensus rather than the original Ripple's local consensus creates additional problems

such as co-ordination of UNL's, incentives for such co-ordination, and unclear incentives for validating nodes to validate correctly or honestly. Finally the requirement of global consensus creates scalability concerns, privacy concerns, and jurisdictional concerns.

A key question that should be answered in future work is if the goals of the Ripple system need global consensus at all? If global consensus can be avoided, or at least its use minimized, many of these issues may go away.

References

- [1] Gavin Andresen. March 2013 chain fork post-mortem. *Bitcoin Improvement Proposals*, 2013. <https://github.com/bitcoin/bips/blob/1647c7e6d3e3866454b803fd874df2d0c4991f4a/bip-0050.mediawiki>.
- [2] Vitalik Buterin. Slasher: A punitive proof-of-stake algorithm. *Ethereum Blog*, January 2014. <https://blog.ethereum.org/2014/01/15/slasher-a-punitive-proof-of-stake-algorithm/>.
- [3] Vitalik Buterin. <https://bitcoinmagazine.com/3506/introducing-ripple/>, 2013. <https://bitcoinmagazine.com/3506/introducing-ripple/>.
- [4] Luke Dashjr. M-of-n standard transactions (low sigop). *Bitcoin Improvement Proposals*, 2012. <https://github.com/bitcoin/bips/blob/1647c7e6d3e3866454b803fd874df2d0c4991f4a/bip-0019.mediawiki>.
- [5] Arthur Britto David Schwartz, Noah Youngs. The ripple protocol consensus algorithm. *Ripple Labs Inc.*, 2014. http://ripple.com/files/ripple_consensus_whitepaper.pdf.
- [6] Stan Higgins. 8 million vericoin hack prompts hard fork to recover funds. *CoinDesk*, July 2014. [Online; accessed 25-April-2015].
- [7] Ripple Labs Inc. Multisign, 2015. [Online; accessed 24-April-2015].
- [8] Ripple Labs Inc. Unique node list, 2015. [Online; accessed 24-April-2015].
- [9] Gregory Maxwell. Proof-of-treachery. #bitcoin-dev IRC channel, July 2012. <http://bitcoinstats.com/irc/bitcoin-dev/logs/2012/07/18#11342572176>.
- [10] Gregory Maxwell. Re: Wtf happened to ripple?, 2013. [Online; accessed 24-April-2015].
- [11] Andrew Miller. Re: Ripple whitepaper questions, 2014. [Online; accessed 24-April-2015].
- [12] Andrew Poelstra. On stake and consensus, March 2015. <https://download.wpsoftware.net/bitcoin/pos.pdf>.

- [13] Bruce Schneier. New nsa leak shows mitm attacks against major internet services. *Schneier on Security*, September 2013. https://www.schneier.com/blog/archives/2013/09/new_nsa_leak_sh.html.
- [14] William Suberg. Ripple will require verified user id to use its compliant network. *Cointelegraph*, 2015. [Online; accessed 25-April-2015].
- [15] Peter Todd. Inflation-proofing via fraud notices. Bitcointalk Forum - Development & Technical Discussion, January 2013. <https://bitcointalk.org/index.php?topic=137933.msg1469812#msg1469812>.
- [16] Wladimir J. van der Laan. Psa: Please sign your git commits. Bitcoin-development mailing list, May 2014. <http://www.mail-archive.com/bitcoin-development@lists.sourceforge.net/msg05467.html>.